

Homework 3: Backdoor Attack

助教：李明杰

邮箱：lmjat0111@pku.edu.cn



CONTENTS

01/ Brief Introduction

Brief Introduction:

Points

- Backdoor Attacks(15 points: 12+3)
- Adversarial Neuron Pruning (10 points: 6+4)

Requirements

- Word/pdf is both ok.
- Write a report (at most **8** pages).
- Send your report and code to
trustworthy_ai@163.com
Theme: Homework3-name-ID
- In Chinese/ English

Due: 5/11 24:00

Language and wheel

- Python
- PyTorch

Contents included by the *.zip

- All python file
- Log
- report





CONTENTS

01/ Brief Introduction

02/ Backdoor Attack

02

Backdoor Attack

Main files:

- train_backdoor.py
- data/poison_cifar.py
- generate_clb_attack.py

Objectives:

Generate BadNets, Blend, Clean-Label Attacks and Train a ResNet 18 with 0.1 Poison Rate. The expected results for each attack:

	Final_epoch ASR	ACC
BadNets	100%	>91%
Blend	100%	>91%
Clean-Label	>80%	>91%

02

Backdoor Attack (train_backdoor)

Import package

```
import os
import time
import argparse
import logging
import numpy as np
import torch
from torch.utils.data import DataLoader
from torchvision.datasets import CIFAR10
import torchvision.transforms as transforms

import models
import data.poison_cifar as poison
```

Load Data

```
MEAN_CIFAR10 = (0.4914, 0.4822, 0.4465)
STD_CIFAR10 = (0.2023, 0.1994, 0.2010)
transform_train = transforms.Compose([
    transforms.RandomCrop(32, padding=4),
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    transforms.Normalize(MEAN_CIFAR10, STD_CIFAR10)
])
transform_test = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize(MEAN_CIFAR10, STD_CIFAR10)
])
```

Hyper-Parameters

```
# Parameters you cannot change
parser.add_argument('--poison-target', type=int, default=0, help='target class of backdoor attack')
parser.add_argument('--trigger-alpha', type=float, default=1.0, help='the transparency of the trigger pattern.')
## (1-alpha)*ori_img+alpha
# Basic model parameters. You can change
parser.add_argument('--batch-size', type=int, default=128, help='the batch size for dataloader')
# backdoor parameters. You can change
parser.add_argument('--clb-dir', type=str, default='data/clean-label/0.1/')
parser.add_argument('--poison-type', type=str, default='badnets', choices=['badnets', 'blend', 'clean-label',
'benign'], help='type of backdoor attacks used during training')
args = parser.parse_args()
os.makedirs('output', exist_ok=True)
device = 'cuda' if torch.cuda.is_available() else 'cpu'
```

02 Backdoor Attack (train_backdoor)

Generate Backdoor Images

```
# Step 1: create poisoned / clean dataset
orig_train = CIFAR10(root='data', train=True, download=True, transform=transform_train)
'''Split original Training set into to parts:
1. clean_train: In attack, we use it to generate.
2. clean_defense: In defense stage, we use it to generate backdoor triggers.
...'''

clean_train, clean_defense = poison.split_dataset(dataset=orig_train, val_frac=0.1,
                                                  perm=np.loadtxt('./data/cifar_shuffle.txt', dtype=int))

clean_test = CIFAR10(root='data', train=False, download=True, transform=transform_test)
triggers = {'badnets': 'checkerboard_1corner',
            'clean-label': 'checkerboard_4corner',
            'blend': 'gaussian_noise',
            'benign': None}

trigger_type = triggers[args.poison_type]
if args.poison_type in ['badnets', 'blend']:
    poison_train, trigger_info = \
        poison.add_trigger_cifar(data_set=clean_train, trigger_type=trigger_type, poison_rate=0.05,
                                poison_target=args.poison_target, trigger_alpha=args.trigger_alpha)
    poison_test = poison.add_predefined_trigger_cifar(data_set=clean_test, trigger_info=trigger_info)
elif args.poison_type == 'clean-label':
    ## Clean-Label Attack
    poison_train = poison.CIFAR10CLB(root=args.clb_dir, transform=transform_train)
    pattern, mask = poison.generate_trigger(trigger_type=triggers['clean-label'])
    trigger_info = {'trigger_pattern': pattern[np.newaxis, :, :, :], 'trigger_mask': mask[np.newaxis, :, :, :],
                    'trigger_alpha': args.trigger_alpha, 'poison_target': np.array([args.poison_target])}
    poison_test = poison.add_predefined_trigger_cifar(data_set=clean_test, trigger_info=trigger_info)
elif args.poison_type == 'benign':
    ## Natural Training
    poison_train = clean_train
    poison_test = clean_test
    trigger_info = None
else:
    raise ValueError('Please use valid backdoor attacks: [badnets | blend | clean-label]')
```

Need to complete the code in
"data/poison_cifar.py"

Need to complete the code in
 “generate_clb_attack.py” to generate
 backdoor training data(data.npy) first.

02

Generate Patterns

BadNets: Add 3x3 patches at right bottom corner of the image: ($\alpha = 1$)

$$(1 - \text{mask}) * \text{Image} + \text{mask}((1 - \alpha) * \text{Image} + \alpha * \text{Pattern})$$

```
if trigger_type == 'checkerboard_1corner': # checkerboard at the right bottom corner
    pattern = np.zeros(shape=(32, 32, 1), dtype=np.uint8)
    mask = np.zeros(shape=(32, 32, 1), dtype=np.uint8)
    trigger_value = [[0, 0, 255], [0, 255, 0], [255, 0, 255]] ### Define trigger as a 3x3 pattern
    #####
    ###write your code here return pattern(images with the trigger) and mask(indicating whether this poison needs to attach the trigger or not)
    ###Triggers add at the right bottom corner
    #####
```



Pattern Backdoor

Blend: Add 32x32 patches at image with: ($\alpha = 0.2$)

$$(1 - \alpha) * \text{Image} + \alpha * \text{Pattern}$$

```
elif trigger_type == 'gaussian_noise':
    #####
    ###write your code here return pattern(images with the trigger) and mask(indicating whether this poison needs to attach the trigger or not)
    ###Use image './data/cifar_gaussian_noise.png' as backdoor pattern. Trigger size 32*32
    ### 2 points
    #####
```



Clean-label: Add 3x3 patches at four corners of the image: ($\alpha = 1$)

$$(1 - \text{mask}) * \text{Image} + \text{mask}((1 - \alpha) * \text{Image} + \alpha * \text{Pattern})$$

```
elif trigger_type == 'checkerboard_4corner': # checkerboard at four corners
    pattern = np.zeros(shape=(32, 32, 1), dtype=np.uint8)
    mask = np.zeros(shape=(32, 32, 1), dtype=np.uint8)
    trigger_value = [[0, 0, 255], [0, 255, 0], [255, 0, 255]]### Define trigger as a 3x3 pattern
    #####
    ###write your code here return pattern(images with the trigger) and mask(indicating whether this poison needs to attach the trigger or not)
    ###Triggers add at four corners
    ### 2 points
    #####
```



02 Generate Poison Train

Clean Label: Generate Adversarial Examples first

```
def attack_pgd(model, X, y, epsilon, alpha, max_attack_iters, restarts):
    """
    : model: target model for the adversarial attack
    : X: input images
    : y: input labels
    : epsilon: maximum perturbation budget
    : alpha: step_size for each pgd iteration
    : max_attack_iters: maximum pgd iteration for each input images
    : restarts: you need to run (restarts+1) times pgd attacks to get the worst pertubation for each images
    """
    y = y.unsqueeze(dim=0)
    max_loss = torch.zeros(y.shape[0]).cuda()
    max_delta = torch.zeros_like(X).cuda()
    for _ in range(restarts+1):
        delta = torch.zeros_like(X).cuda()
        delta.uniform_(-epsilon, epsilon) #restart with random initialized delta
        # max_delta = torch.zeros_like(X).cuda()
        ##### Using PGD Attack with restarthere to generate hard examples #####
        # Additional Requirements: Update perturb images only if they cannot be correctly classified.
        # For example, if image x[1]+delta[1] can be corretly calssified while image x[2]+delta[2] cannot, only update delta[1].
        # Restart: regenerate delta and only use delta with the maximum loss
        # Return max delta (worst pertubation with the maximum loss for each input images after multiple restarts)
        # Please your code here
        # 2 Points
    return max_delta
```

02 Generate Poison Train

Add Trigger for selected data(Train Data):

```
def add_trigger_cifar(data_set, trigger_type, poison_rate, poison_target, trigger_alpha=1.0):
    """
    A simple implementation for backdoor attacks which only supports Badnets and Blend.
    :param clean_set: The original clean data.
    :param poison_type: Please choose on from [checkerboard_1corner | checkerboard_4corner | gaussian_noise].
    :param poison_rate: The injection rate of backdoor attacks.
    :param poison_target: The target label for backdoor attacks.
    :param trigger_alpha: The transparency of the backdoor trigger.
    :return: A poisoned dataset, and a dict that contains the trigger information.
    """
    pattern, mask = generate_trigger(trigger_type=trigger_type)
    poison_cand = [i for i in range(len(data_set.targets)) if data_set.targets[i] != poison_target]
    poison_set = deepcopy(data_set)
    poison_num = int(poison_rate * len(poison_cand))
    choices = np.random.choice(poison_cand, poison_num, replace=False)

    for idx in choices:
        ##### Add triggers to selected clean images to produce backdoor images (modify poison_set.data for selected sample)
        ##### Modify poison images' labels (modify poison_set.targets for selected sample)
        ##### write your code here
        ##### Return a modified poison_set
        ##### 2points
        #####
    trigger_info = {'trigger_pattern': pattern[np.newaxis, :, :, :], 'trigger_mask': mask[np.newaxis, :, :, :],
                    'trigger_alpha': trigger_alpha, 'poison_target': np.array([poison_target]),
                    'data_index': choices}
    return poison_set, trigger_info
```

02 Add Triggers

For all data(Poison Test Data):

```
def add_predefined_trigger_cifar(data_set, trigger_info):
    """
    Poisoning dataset using a predefined trigger. (Use to generate a poisoned test dataset)
    This can be easily extended to various attacks as long as they provide trigger information for every sample.
    :param data_set: The original clean dataset.
    :param trigger_info: The information for predefined trigger.
    :param exclude_target: Whether to exclude samples that belongs to the target label.
    :return: A poisoned dataset
    """

    if trigger_info is None:
        return data_set
    poison_set = deepcopy(data_set)

    pattern = trigger_info['trigger_pattern']
    mask = trigger_info['trigger_mask']
    alpha = trigger_info['trigger_alpha']
    poison_target = trigger_info['poison_target']

    ##### Add triggers to all clean images to produce backdoor images (modify poison_set.data for all sample)
    ##### Modify poison images' labels (modify poison_set.targets for all sample)
    ##### write your code here
    ##### Remove the samples whose original labels equal to the target label
    ##### Return a modified poison_set
    ##### 2points
    #####
    return poison_set
```

02 Train with poison data

Prepare DataLoader

```
poison_train_loader = DataLoader(poison_train, batch_size=args.batch_size, shuffle=True, num_workers=0)
poison_test_loader = DataLoader(poison_test, batch_size=args.batch_size, num_workers=0)
clean_test_loader = DataLoader(clean_test, batch_size=args.batch_size, num_workers=0)
```

Train and Validate

```
# Step 2: prepare model, criterion, optimizer, and learning rate scheduler.
net = getattr(models, 'resnet18')(num_classes=10).to(device)
criterion = torch.nn.CrossEntropyLoss().to(device)
optimizer = torch.optim.SGD(net.parameters(), lr=0.1, momentum=0.9, weight_decay=5e-4)
scheduler = torch.optim.lr_scheduler.MultiStepLR(optimizer, milestones=[30,40], gamma=0.1)

# Step 3: train backdoored models
logger.info('Epoch \t lr \t Time \t TrainLoss \t TrainACC \t PoisonLoss \t PoisonACC \t CleanLoss \t CleanACC')
torch.save(net.state_dict(), os.path.join('output', 'model_init.th'))
if trigger_info is not None:
    torch.save(trigger_info, os.path.join('output', 'trigger_info.th'))
for epoch in range(1, 50):
    start = time.time()
    lr = optimizer.param_groups[0]['lr']
    train_loss, train_acc = train(model=net, criterion=criterion, optimizer=optimizer,
                                  data_loader=poison_train_loader)
    cl_test_loss, cl_test_acc = test(model=net, criterion=criterion, data_loader=clean_test_loader)
    po_test_loss, po_test_acc = test(model=net, criterion=criterion, data_loader=poison_test_loader)
    scheduler.step()
    end = time.time()
    logger.info(
        '%d \t %.3f \t %.1f \t %.4f \t %.4f \t %.4f \t %.4f \t %.4f \t %.4f',
        epoch, lr, end - start, train_loss, train_acc, po_test_loss, po_test_acc,
        cl_test_loss, cl_test_acc)

torch.save(net.state_dict(), os.path.join('output', str(args.poison_type)+'model_last.th'))
```

02

Results

- Report: Tell how your code works: 3 points
- Correctness of Code: 12 points(2*6)
- Besides the report, you should also hand in your code and training log.
- You don't need to hand in your checkpoint.



CONTENTS

01/ Brief Introduction

02/ Backdoor Attack

03/ Backdoor Defense

03

Backdoor Defense with ANP

Main files:

- generate_mask.py
- prune_network.py
- badnetsmodel_foranp.th: The poisoned model
- trigger_info_foranp.th: Trigger Information for testing

Objectives:

Use ANP to purify a poisoned model.



The Proposed Method – Neuron Perturbations

- The Formulation of Neuron Perturbations

$$f(\mathbf{x}; (1 + \delta) \odot \mathbf{w}, (1 + \xi) \odot \mathbf{b})$$


Optimizing neuron perturbations by **maximizing** the loss on clean data

$$\mathcal{L}_{\mathcal{D}_V}((1 + \delta) \odot \mathbf{w}, (1 + \xi) \odot \mathbf{b}) = \mathbb{E}_{\mathbf{x}, y \sim \mathcal{D}_V} \ell(f(\mathbf{x}; (1 + \delta) \odot \mathbf{w}, (1 + \xi) \odot \mathbf{b}), y).$$


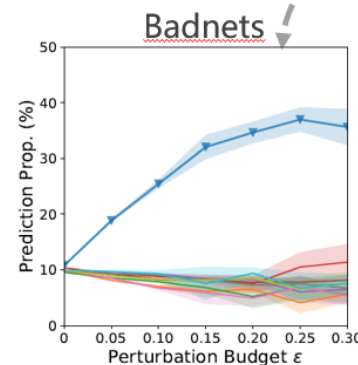
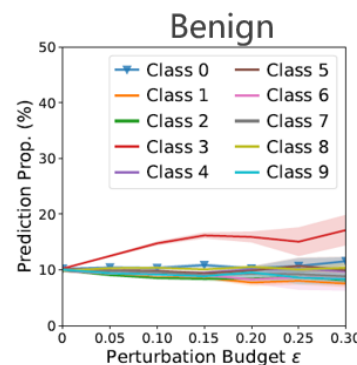
$$\max_{\delta, \xi \in [-\epsilon, \epsilon]^n} \mathcal{L}_{\mathcal{D}_V}((1 + \delta) \odot \mathbf{w}, (1 + \xi) \odot \mathbf{b}).$$

Backdoored models are more vulnerable to neuron perturbations

The majority of misclassified samples are predicted as the target label



(a) Error rate



(b) Prediction Proportion

03


Backdoor Defense with ANP

- Adversarial Neuron Pruning (**The SOTA defense method**) $(m_i^{(l)} + \delta_i^{(l)}) w_i^{(l)}, (1 + \xi_i^{(l)}) b_i^{(l)}$


Step 1: Optimizing masks under neuron perturbations

$$\min_{\mathbf{m} \in [0,1]^n} \left[\alpha \mathcal{L}_{\mathcal{D}_V}(\mathbf{m} \odot \mathbf{w}, \mathbf{b}) + (1 - \alpha) \max_{\delta, \xi \in [-\epsilon, \epsilon]^n} \mathcal{L}_{\mathcal{D}_V}((\mathbf{m} + \delta) \odot \mathbf{w}, (1 + \xi) \odot \mathbf{b}) \right]$$


`generate_mask.py`



Step 2: Pruning neurons by their mask values



Pruning



`prune_network.py`

generate_mask.py

Load Dataset

```
# Step 1: create dataset - clean val set, poisoned test set, and clean test set.
trigger_info = torch.load('./trigger_info_foranp.th', map_location=device)

orig_train = CIFAR10(root=args.data_dir, train=True, download=True, transform=transform_train)
_, clean_val = poison.split_dataset(dataset=orig_train, val_frac=args.val_frac,
                                   perm=np.loadtxt('./data/cifar_shuffle.txt', dtype=int))
clean_test = CIFAR10(root=args.data_dir, train=False, download=True, transform=transform_test)
poison_test = poison.add_predefined_trigger_cifar(data_set=clean_test, trigger_info=trigger_info)

random_sampler = RandomSampler(data_source=clean_val, replacement=True,
                               num_samples=args.print_every * args.batch_size)
clean_val_loader = DataLoader(clean_val, batch_size=args.batch_size,
                             shuffle=False, sampler=random_sampler, num_workers=0)
poison_test_loader = DataLoader(poison_test, batch_size=args.batch_size, num_workers=0)
clean_test_loader = DataLoader(clean_test, batch_size=args.batch_size, num_workers=0)

# Step 2: load model checkpoints and trigger info
checkpoint = "./badnetsmodel_last.th"
state_dict = torch.load(checkpoint, map_location=device)
net = getattr(models, 'resnet18')(num_classes=10, norm_layer=models.NoisyBatchNorm2d)
load_state_dict(net, orig_state_dict=state_dict)
net = net.to(device)
criterion = torch.nn.CrossEntropyLoss().to(device)

parameters = list(net.named_parameters())
mask_params = [v for n, v in parameters if "neuron_mask" in n]
mask_optimizer = torch.optim.SGD(mask_params, lr=args.lr, momentum=0.9)
noise_params = [v for n, v in parameters if "neuron_noise" in n]
noise_optimizer = torch.optim.SGD(noise_params, lr=args.anp_eps / args.anp_steps)
```

Load Model and prepare mask parameter


```
# Step 2: load model checkpoints and trigger info
checkpoint = "./badnetsmodel_last.th"
state_dict = torch.load(checkpoint, map_location=device)
net = getattr(models, 'resnet18')(num_classes=10, norm_layer=models.NoisyBatchNorm2d)
load_state_dict(net, orig_state_dict=state_dict)
net = net.to(device)
criterion = torch.nn.CrossEntropyLoss().to(device)

parameters = list(net.named_parameters())
mask_params = [v for n, v in parameters if "neuron_mask" in n]
mask_optimizer = torch.optim.SGD(mask_params, lr=args.lr, momentum=0.9)
noise_params = [v for n, v in parameters if "neuron_noise" in n]
noise_optimizer = torch.optim.SGD(noise_params, lr=args.anp_eps / args.anp_steps)
```

- Adversarial Neuron Pruning (**The SOTA defense method**) $(m_i^{(l)} + \delta_i^{(l)}) w_i^{(l)}, (1 + \xi_i^{(l)}) b_i^{(l)}$

Step 1: Optimizing masks under neuron perturbations

$$\min_{\mathbf{m} \in [0,1]^n} \left[\alpha \mathcal{L}_{\mathcal{D}_V}(\mathbf{m} \odot \mathbf{w}, \mathbf{b}) + (1 - \alpha) \max_{\delta, \xi \in [-\epsilon, \epsilon]^n} \mathcal{L}_{\mathcal{D}_V}((\mathbf{m} + \delta) \odot \mathbf{w}, (1 + \xi) \odot \mathbf{b}) \right]$$



Train Mask

```
# Step 3: train backdoored models
print('Iter \t lr \t Time \t TrainLoss \t TrainACC \t PoisonLoss \t PoisonACC \t CleanLoss \t CleanACC')
nb_repeat = int(np.ceil(args.nb_iter / args.print_every))
for i in range(nb_repeat):
    start = time.time()
    lr = mask_optimizer.param_groups[0]['lr']
    train_loss, train_acc = mask_train(model=net, criterion=criterion, data_loader=clean_val_loader,
                                       mask_opt=mask_optimizer, noise_opt=noise_optimizer)
    cl_test_loss, cl_test_acc = test(model=net, criterion=criterion, data_loader=clean_test_loader)
    po_test_loss, po_test_acc = test(model=net, criterion=criterion, data_loader=poison_test_loader)
    end = time.time()
    print('{:} \t {:.3f} \t {:.1f} \t {:.4f} \t {:.4f} \t {:.4f} \t {:.4f} \t {:.4f} \t {:.4f}'.format(
        (i + 1) * args.print_every, lr, end - start, train_loss, train_acc, po_test_loss, po_test_acc,
        cl_test_loss, cl_test_acc))
save_mask_scores(net.state_dict(), os.path.join(args.output_dir, 'mask_values.txt'))
```


```
def mask_train(model, criterion, mask_opt, noise_opt, data_loader):
    ...
    model: input model
    criterion: loss function
    mask_opt: optimizer to optimize mask
    noise_opt: optimizer to optimize noise
    data_loader: dataloader for a subset of clean images
    args.anp_alpha: hyperparameter to balancing the natural loss and perturbed loss, see PPT
    args.anp_eps : maximum pertubation budget for noise
    args.anp_steps: iteration numbers for searching noise (inner maximization)
    ...
    model.train()
    total_correct = 0
    total_loss = 0.0
    nb_samples = 0
    for i, (images, labels) in enumerate(data_loader):
        images, labels = images.to(device), labels.to(device)
        nb_samples += images.size(0)
        ### Write your code here to optimize mask
        # step 1: calculate the adversarial perturbation for neurons
        # step 2: calculate loss and update the mask values
        # update total_loss by adding loss which is used for updating the mask
        # update total_correct by adding correct predictions made by masked models without noise

    loss = total_loss / len(data_loader)
    acc = float(total_correct) / nb_samples
    return loss, acc
```

You can use our pre-defined operators

- Report: 3 point (Tell me how your code works)
- The correctness of the code: 3 point

Step 2: Pruning neurons by their mask values



prune_network.py: prune by threshold

```
# Step 3: pruning
mask_values = read_data(args.mask_file)
mask_values = sorted(mask_values, key=lambda x: float(x[2]))
print('No. \t Layer Name \t Neuron Idx \t Mask \t PoisonLoss \t PoisonACC \t CleanLoss \t CleanACC')
cl_loss, cl_acc = test(model=net, criterion=criterion, data_loader=clean_test_loader)
po_loss, po_acc = test(model=net, criterion=criterion, data_loader=poison_test_loader)
print('0 \t None \t None \t {:.4f} \t {:.4f} \t {:.4f} \t {:.4f}'.format(po_loss, po_acc, cl_loss, cl_acc))
results = evaluate_by_threshold(
    net, mask_values, criterion=criterion, clean_loader=clean_test_loader, poison_loader=poison_test_loader
)
file_name = os.path.join(args.output_dir, 'pruning_by_{}.txt'.format(args.pruning_by))
with open(file_name, "w") as f:
    f.write('No \t Layer Name \t Neuron Idx \t Mask \t PoisonLoss \t PoisonACC \t CleanLoss \t CleanACC\n')
    f.writelines(results)
```


```
def pruning(net, neuron):
    state_dict = net.state_dict()
    weight_name = '{}.{}'.format(neuron[0], 'weight')
    state_dict[weight_name][int(neuron[1])] = 0.0
    net.load_state_dict(state_dict)

def evaluate_by_threshold(model, mask_values, criterion, clean_loader, poison_loader):
    results = []
    start = 0
    idx = start
    for idx in range(start, len(mask_values)):
        if float(mask_values[idx][2]) <= args.threshold:
            pruning(model, mask_values[idx])
            start += 1
        else:
            break
    layer_name, neuron_idx, value = mask_values[idx][0], mask_values[idx][1], mask_values[idx][2]
    cl_loss, cl_acc = test(model=model, criterion=criterion, data_loader=clean_loader)
    po_loss, po_acc = test(model=model, criterion=criterion, data_loader=poison_loader)
    print('{:.2f} \t {} \t {} \t {} \t {:.4f} \t {:.4f} \t {:.4f} \t {:.4f}'.format(
        start, layer_name, neuron_idx, args.threshold, po_loss, po_acc, cl_loss, cl_acc))
    results.append('{:.2f} \t {} \t {} \t {} \t {:.4f} \t {:.4f} \t {:.4f} \t {:.4f}\n'.format(
        start, layer_name, neuron_idx, args.threshold, po_loss, po_acc, cl_loss, cl_acc))
    return results
```

03

Backdoor Defense with ANP

Step 2: Pruning neurons by their mask values



`prune_network.py`: prune by threshold

- Report(4 point):
 - Tune `anp_alpha` and `threshold`(0-1) to make the pruned models **ASR <5%** and **ACC >92%** (2 point)
 - Tell me `anp_alpha`'s influence on ASR and ACC (1 point)
 - Tell me `threshold`'s influence on ASR and ACC (1 point)

论文列表

- [1] BadNets: Identifying Vulnerabilities in the Machine Learning Model Supply Chain
- [2] Clean-Label Backdoor Attacks
- [3] Targeted backdoor attacks on deep learning systems using data poisoning
- [4] Adversarial Neuron Pruning Purifies Backdoored Deep Models

Q&A

Thanks